

# Bare-Metal Forth

*Implementing and Programming Forth on Microcontrollers*

---

## Preface — Scope, Philosophy, and Audience

This work is **both a course and a book**.

### Audience

- Readers **new to Forth** (no prior Forth experience assumed)
- Embedded developers curious about **language–machine intimacy**
- Learners who want to understand *why* Forth looks the way it does

### What this book teaches

#### 1. The philosophy of Forth

- Smallness, extensibility, and trust in the programmer
- Interactive development and the REPL as a thinking tool
- Words as vocabulary; vocabularies as worldviews

#### 2. eForth style and lineage

- Minimal, portable kernels
- Clear separation of a tiny machine-dependent core and portable ideas
- Learning by *rebuilding*, not by consuming frameworks

#### 3. Implementation and use, side by side

- How a Forth system is built on bare metal
- How the same system is *used* interactively

### What this book deliberately avoids

- Tying explanations to **one specific threading model** (ITC/STC/NTC)
- Treating Forth as a historical curiosity or niche language
- Hiding mechanisms behind opaque tooling or HALs

Threading models, memory layouts, and optimizations are treated as **design choices**, not dogma.

## Historical remarks and references

- Historical notes (figForth, eForth, AmForth, Mecrisp, etc.) are included **where they clarify a design decision**
- References and corrections are woven into the narrative, not isolated in footnotes
- Code corrections and alternatives are explained as *consequences of constraints*, not as errors

## Structure

- Each chapter exists as:
  - a **teaching text** (this manuscript)
  - a **self-contained assembly stage** that can be built and studied independently

The goal is not to teach “one true Forth”, but to teach **how Forth comes into being on real hardware** — and how that shapes the way one thinks.

---

# Lesson 1 — Reset, Memory, and the Empty Machine

*Before there is a language, there is a machine. Before there is a machine, there is reset.*

## Purpose of this lesson

This lesson introduces the **absolute minimum** required to bring a microcontroller from power-on into a *known, controlled state*. No Forth is executed yet. No words are defined. We deliberately stay *before* the language.

By the end of this lesson, the reader should:

- understand what happens on reset
- know why memory must be cleared deliberately
- distinguish **hardware return stack** from **software data stack**
- see memory as a *finite territory*, not an abstraction

This lesson is about **discipline**, not features.

---

## 1. Reset as a semantic boundary

On a microcontroller, reset is not just a restart. It is the only moment where:

- the hardware promises a defined entry point
- the programmer has full authority over initial conditions

Everything that follows inherits the quality of this moment.

On the ATtiny85, reset transfers control to address **0x0000**, the **reset vector**.

---

## 2. Interrupt vector table (minimal form)

In this lesson, the vector table is intentionally minimal:

- the reset vector jumps to our reset handler
- all other vectors immediately return (**RETI**)

This achieves two things:

1. the machine is safe from unintended interrupts
  2. the reader sees clearly that *nothing happens unless we allow it*
- 

## 3. Capturing the reset cause

The MCU provides a register (**MCUSR**) that records *why* the reset occurred:

- power-on
- watchdog
- brown-out
- external reset

We read this register once, at reset, and then clear it.

This teaches an important rule:

**Information that is not captured immediately may be lost forever.**

Even if we do not use the value yet, the act of capturing it establishes good practice.

---

## 4. Clearing SRAM — rejecting randomness

SRAM contents after reset are **undefined**.

Any system that assumes otherwise is relying on coincidence.

We therefore explicitly clear the entire usable SRAM region:

- starting at address **0x0060**
- ending at the top of SRAM

This is not an optimization. It is a declaration:

*There will be no ghosts from previous executions.*

---

## 5. Two stacks, two meanings

At this point we introduce the first crucial distinction:

- **Return stack**
  - implemented by the AVR hardware stack pointer (**SP**)
  - used implicitly by **CALL / RET**
- **Data stack**
  - implemented in software
  - explicitly managed by the program

In this system:

- the return stack is initialized by writing to **SP**
- the data stack pointer is stored in a register pair (**Y**)

This separation is the foundation of Forth.

---

## 6. Memory map as geography

Rather than treating RAM as a pool of variables, we divide it into regions:

- user area
- input buffer
- free workspace
- data stack
- return stack

Each region has:

- a start
- an end
- a purpose

The exact addresses matter less than the *act of choosing them deliberately*.

---

## 7. Proof of life (optional but human)

A blinking LED is not required by the machine.

It is required by the human.

A single GPIO pin configured as output and toggled in a loop serves as:

- confirmation that reset worked
  - confirmation that timing exists
  - confirmation that control is ours
- 

# Exercises

## Exercise 1 — Memory without clearing

Comment out the SRAM clearing loop.

- Flash the program.
- Reset the device multiple times.
- Observe whether behavior remains deterministic.

**Question:** What assumptions did the original version refuse to make?

---

## Exercise 2 — Moving the data stack

Change the initial data stack pointer to a lower address.

- How much space does the stack really need *at this stage*?
  - What happens if it overlaps another region?
- 

## Exercise 3 — Breaking the return stack

Initialize the hardware stack pointer incorrectly.

- What fails first?
  - Does the failure look predictable or chaotic?
- 

## Exercise 4 — Draw the memory map

On paper, draw the SRAM layout:

---

- before reset
- after reset

Label each region and its purpose.

This drawing will be reused throughout the book.

---

## Reflection

This lesson contains no Forth code.

And yet, without it, no Forth system can exist.

Before a language can be interactive, it must be *stable*.

---